


CUDA

By: Guillermo Andres De Mendoza Corrales



CPU vs GPU



MakeAGIF.com

- Central processing unit
- **Versatile**
- 1 to 30 cores
- Low latency GHz
- Serial



- Graph processing unit
- **Only some task (vectors)**
- 100s to 10000s cores
- High throughput (MHz)
- Parallel



Architecture comparison



ALU: Arithmetic Logic Unit

Control Unit:

- Search instructions
- Decode instructions
- Control executions
- Manage flows *by buss
- Synchronize

Cache:

- Contain instructions
- Data access
- L1 -> L2 -> L3 -> RAM




GPU



CUDA

CUDA (Compute Unified Device Architecture) is an architecture for parallel computing, developed by **NVIDIA** in 2007, that leverages the power of its GPUs for general-purpose processing (**not only graphical**). It uses a hierarchy of grids, thread blocks, and threads to execute tasks in parallel, significantly speeding up compute-intensive application.



CC : Compute Capability

Define the allowed functions and characteristics that the GPU can perform.

Each CUDA version require a minimum CC value

Some advance functions require also a minimum CC value

- Tensor cores
- Ray tracing
- Acceleration FP16
- Some libraries

Capacidad de cálculo	Centro de datos	Estación de trabajo/Consumidor	Supersónico
12.0	NVIDIA RTX PRO 6000 Blackwell Server Edition	NVIDIA RTX PRO 6000 Blackwell Edición para estación de trabajo NVIDIA RTX PRO 6000 Blackwell Edición para estación de trabajo Max-Q NVIDIA RTX PRO 5000 Blackwell NVIDIA RTX PRO 4500 Blackwell NVIDIA RTX PRO 4000 Blackwell NVIDIA RTX PRO 4000 Blackwell Edición SFF NVIDIA RTX PRO 2000 Blackwell GeForce RTX 5090 GeForce RTX 5080 GeForce RTX 5070 Ti GeForce RTX 5070 GeForce RTX 5060 Ti GeForce RTX 5060 GeForce RTX 5050	
11.0			Jetson T5000 Jetson T4000
10.3	NVIDIA GB300 NVIDIA B300		
10.0	NVIDIA GB200 NVIDIA B200		
9.0	NVIDIA GH200 NVIDIA H200 NVIDIA H100	NVIDIA RTX 6000 Ada NVIDIA RTX 5000 Ada NVIDIA RTX 4000 Ada NVIDIA RTX 4000 Ada NVIDIA RTX 4000 SFF Ada NVIDIA RTX 2000 Ada GeForce RTX 4090 GeForce RTX 4080 GeForce RTX 4070 Ti GeForce RTX 4070 GeForce RTX 4060 Ti GeForce RTX 4060 GeForce RTX 4050	
8.9	NVIDIA L4 NVIDIA L40 NVIDIA L40S		
8.7			Jetson AGX Orin Jetson Orin NX Jetson Orin Nano
8.6	NVIDIA A40 NVIDIA A10 NVIDIA A16 NVIDIA A2	NVIDIA RTX A6000 NVIDIA RTX A5000 NVIDIA RTX A4000 NVIDIA RTX A3000 NVIDIA RTX A2000 GeForce RTX 3090 Ti GeForce RTX 3090 GeForce RTX 3080 Ti GeForce RTX 3080 GeForce RTX 3070 Ti GeForce RTX 3070 GeForce RTX 3060 Ti GeForce RTX 3060 GeForce RTX 3050 Ti	



KERNEL

Each thread execute a copy of the kernel







Imagen: y:110 , x:100

ThreadBlock: 15x15



Grid: (x:7, y:8)



KERNEL

1. CUDA Toolkit

- CUDA compiler nvcc
- CUDA Libraries
- CUDA Runtime
- CUDA Driver APIs
- CUDA Debugging
- CUDA Performance Analysis

Download Installer for Windows 11 x86_64

The base installer is available for download below.

> CUDA Toolkit Installer Download (2.3 GB)

Installation Instructions:

1. Double click cuda_13.0.2_windows.exe
2. Follow on-screen prompts

Additional installation options are detailed [here](#).

Command Prompt



Microsoft Windows [Version 10.0.26100.6899]
(c) Microsoft Corporation. All rights reserved.

```
C:\Users\memoo>nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2025 NVIDIA Corporation
Built on Wed_Aug_20_13:58:20_Pacific_Daylight_Time_2025
Cuda compilation tools, release 13.0, V13.0.88
Build cuda_13.0.r13.0/compiler.36424714_0
```



COLAB IMPLEMENTATION

Set notebook to GPU



Hardware environment

[1]
✓ 4s

```
from numba import cuda
print(cuda.is_available())
cuda.detect()

True
Found 1 CUDA devices
id 0          b'Tesla T4'           [SUPPORTED]
              Compute Capability: 7.5
              PCI Device ID: 4
              PCI Bus ID: 0
              UUID: GPU-4637cb14-ea62-ed4a-90d3-2709af8ff9c3
              Watchdog: Disabled
              FP32/FP64 Performance Ratio: 32
Summary:
  1/1 devices are supported
True
```



GPU Architecture	NVIDIA Turing
NVIDIA Turing Tensor Cores	320
NVIDIA CUDA® Cores	2,560
Single-Precision	8.1 TFLOPS
Mixed-Precision (FP16/FP32)	65 TFLOPS
INT8	130 TOPS
INT4	260 TOPS
GPU Memory	16 GB GDDR6 300 GB/sec
ECC	Yes
Interconnect Bandwidth	32 GB/sec
System Interface	x16 PCIe Gen3
Form Factor	Low-Profile PCIe
Thermal Solution	Passive
Compute APIs	CUDA, NVIDIA TensorRT™, ONNX



Hardware environment

```
!nvidia-smi

Sat Oct 18 02:49:58 2025
+-----+
| NVIDIA-SMI 550.54.15      Driver Version: 550.54.15     CUDA Version: 12.4 |
| Persistence-M | Bus-Id     Disp.A  | Volatile Uncorr. ECC | | | | | |
| GPU  Name      Fan  Temp   Perf  Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
|          |             |          |          |          |          |          |          |
| 0  Tesla T4    N/A   56C   P8    Off  00000000:00:04.0 Off | 0MiB / 15360MiB | 0%       Default |
|          |             |          |          |          |          |          |          |
+-----+
Processes:
| GPU  GI  CI          PID  Type  Process name          GPU Memory |
|          ID  ID          ID          ID          Usage      |
+-----+
| No running processes found
+-----+
```



CODE

Avoid errors

```
from google.colab import files  
import cv2  
import os  
import shutil  
import matplotlib.pyplot as plt
```

```
import numpy as np  
from numba import cuda  
from numba import config  
config.CUDA_ENABLE_PYNVJITLINK = 1 # resuelve bug de versiones
```



Kernel

```
@cuda.jit
def grayscale_kernel_2D(img_color, img_gray):
    """
        Kernel 2D que mapea las coordenadas (i, j) directamente usando
        índices de bloque y de hilo en las dimensiones X e Y.
        i = fila (altura), j = columna (ancho)
    """
    # Obtener las coordenadas 2D (i, j)
    # cuda.grid(2) devuelve (j, i) -> (column, row)
    j, i = cuda.grid(2)

    # Obtener las dimensiones de la imagen de salida (filas, columnas)
    rows, cols = img_gray.shape

    # Chequeo de límites
    # Asegurarse de que las coordenadas estén dentro de los límites de la imagen
    if i < rows and j < cols:
        # Acceder a los componentes de color de la imagen original
        R = img_color[i, j, 0]
        G = img_color[i, j, 1]
        B = img_color[i, j, 2]

        # Cálculo de la intensidad de gris
        # Se mantienen los coeficientes estándar
        gray_value = 0.2989 * R + 0.5870 * G + 0.1140 * B

        # Asignar el valor de gris al array de salida
        img_gray[i, j] = gray_value
```



GPU Caller

```
def image_to_grayscale(img):

    # 'img' entra como uint8 (0-255)
    ROWS, COLUMNS, COLOR = img.shape


    # Transferir la imagen original a la GPU
    d_img_color = cuda.to_device(img)

    # Reservar memoria en GPU para la imagen de salida
    d_img_gray = cuda.device_array((ROWS, COLUMNS), dtype=np.float32)

    # Configuración del GRID
    threadsperblock = (16, 16)
    blockspergrid_x = (COLUMNS + threadsperblock[0] - 1) // threadsperblock[0]
    blockspergrid_y = (ROWS + threadsperblock[1] - 1) // threadsperblock[1]
    blockspergrid = (blockspergrid_x, blockspergrid_y)

    # Lanzamiento
    grayscale_kernel_2D[blockspergrid, threadsperblock](d_img_color, d_img_gray)

    # Recuperar (copy_to_host ya sincroniza implícitamente)
    return d_img_gray.copy_to_host()
```



Read Image


```
uploaded = files.upload()

img_color_name = list(uploaded.keys())[0]
img_color = cv2.imread(img_color_name)
img_color_cv_rgb = cv2.cvtColor(img_color, cv2.COLOR_BGR2RGB)
img_color_np = img_color_cv_rgb.astype(np.float32) / 255.0
H, W, C = img_color_np.shape
print(f"Imagen cargada como NumPy array. Forma (H, W, C): {H}x{W}x{C}")

plt.imshow(img_color_np)
plt.show()
```

Elegir archivos

pikachu.png
pikachu.png(image/png) - 462351 bytes, last modified: 5/2/2026 - 100% done
Saving pikachu.png to pikachu.png
Imagen cargada como NumPy array. Forma (H, W, C): 948x812x3



Logic

```
img_gray = image_to_grayscale(img_color_np)

plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
# img_np debe normalizarse si no está en el rango [0, 255]
plt.imshow(img_color_np / 255 if img_color_np.max() > 1 else img_color_np)
plt.title("Imagen Original")
plt.axis('off')

plt.subplot(1, 2, 2)
# El resultado es float32, usamos cmap='gray'
plt.imshow(img_gray, cmap='gray')
plt.title("Imagen en Blanco y Negro (CUDA 2D)")
plt.axis('off')

plt.show()
```

